


Parameter Estimation

(+ a Mean Example)



Estimators

p Unknown distribution or population

θ Parameter of p wish to estimate Mean example μ σ , max, mode

Sample $X^n \stackrel{\text{def}}{=} X_1, X_2, \dots, X_n \sim p$

$X^3 = 5, -2, 6$

$\hat{\theta}$ for estimate

Estimator for θ function $\hat{\theta} : \mathbb{R}^n \rightarrow \mathbb{R}$

Maps X^n to \mathbb{R}

Upon observing X^n , estimate θ as $\hat{\theta}(X^n) \stackrel{\text{def}}{=} \hat{\Theta}$

μ $\hat{\theta}(X^n) :$ $\frac{X_1 + \dots + X_n}{n}$ $\frac{\min\{X_i\} + \max\{X_i\}}{2}$ $X_1 \cdot X_2$

5, -2, 6 $\hat{\Theta} :$

3

2

-10

Observations

Distribution parameter θ

Constant

Mean 3.2

Estimate

$$\hat{\Theta} \stackrel{\text{def}}{=} \hat{\theta}(X^n)$$

Random variable

Ideally close to θ

Once sample X^n drawn

Determines $\hat{\Theta}$

Single value

3.5

Point estimate

vs. interval

[3,4]

Any function is an estimator

Come up with an estimator?

How to

Evaluate its performance?



SPACEX
Elon Musk
Founder of SpaceX



SAMPLE ✕

Apply sample to any parameter ✕

Sample X

Property

X

min

X_{\min}

$$\min_x \{x : p(x) > 0\}$$

max

X_{\max}

$$\max_x \{x : p(x) > 0\}$$

mean

μ

$$\sum_x x \cdot p(x)$$

Simple

If sample is whole population, exact

Sometimes works well

Even for small samples

sample X

sample min

$$\min_i \{X_i\}$$

sample max

$$\max_i \{X_i\}$$


sample mean

$$\frac{1}{n} \sum_{i=1}^n X_i$$


Estimator Evaluation

Parameter may have several estimators


Evaluate quality of estimator for a parameter



Good



Bias



Variance


Bias

$\hat{\theta}$ estimator for θ

Bias of $\hat{\theta}$ is its expected overestimate of θ


$$\text{Bias}_\theta(\hat{\theta}) \stackrel{\text{def}}{=} E(\hat{\theta} - \theta) = \mu_{\hat{\theta}} - \theta$$

$$\rightarrow \text{Bias}(\hat{\theta})$$



Estimator with 0 bias is **unbiased**

$$\mu_{\hat{\theta}} = \theta$$



Bias = Inequality

Variance

$$V(\hat{\Theta}) = E(\hat{\Theta} - \mu_{\hat{\Theta}})^2$$

Unrelated to θ

Ideally 0 bias variance

Typically tradeoff

Mean Example

Unknown distribution or population p

Estimate mean μ

n samples

$$X_1, \dots, X_n \sim p \perp\!\!\!\perp$$

Sample mean

$$\bar{X} \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n X_i$$

Evaluate

Bias

Variance



Weak Law of Large Numbers

Sample Mean - Bias

Sample mean

$$\overline{X} \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n X_i$$

Expectation

$$\begin{aligned} E(\overline{X}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} \sum_{i=1}^n E(X_i) = \mu \end{aligned}$$

Bias


$$\text{Bias}(\overline{X}) = E(\overline{X}) - \mu = \mu - \mu = 0$$

Sample mean is unbiased estimator for distribution mean

Sample Mean - Variance

$$V(\bar{X}) = V\left(\frac{1}{n} \sum_{i=1}^n X_i\right)$$

$$= \frac{1}{n^2} V\left(\sum_{i=1}^n X_i\right)$$



$$= \frac{1}{n^2} \sum_{i=1}^n V(X_i)$$





$$= \frac{1}{n^2} \sum_{i=1}^n \sigma^2$$

$$= \frac{\sigma^2}{n} \qquad \qquad \sigma_{\bar{X}} = \frac{\sigma}{\sqrt{n}}$$

Increases with σ

Decreases with n

Experiments



Mean Squared Error

Single measure for performance of estimator $\hat{\Theta}$ for θ

MSE of $\hat{\Theta}$ is its expected squared distance from θ

$$\text{MSE}_\theta(\hat{\Theta}) \stackrel{\text{def}}{=} E(\hat{\Theta} - \theta)^2 \rightarrow \text{MSE}(\hat{\Theta})$$

Common in science and engineering

Communication

Transportation

Production

Need to re-evaluate?

Relate to bias and variance

Bias-Variance Bromance

$$\text{MSE} = \text{ Bias}^2 + \text{Variance}$$

$$\text{MSE}(\Theta) = E(\Theta - \theta)^2$$


$$= E^2(\Theta - \theta) + V(\Theta - \theta)$$

$$E(\Theta - \theta) \stackrel{\text{def}}{=} \text{Bias}(\Theta)$$

$$V(\Theta - \theta) = V(\Theta)$$

$$= \text{Bias}^2(\Theta) + V(\Theta)$$

$$E(X^2) = E^2(X) + V(X)$$



$$\text{Energy} = \mu^2 + \sigma^2$$

MSE of Sample Mean

$$\text{MSE}_\mu(\bar{X}) = \text{Bias}_\mu^2(\bar{X}) + V(\bar{X}) = \frac{\sigma^2}{n}$$

Increases with σ


Decreases with n

Same estimator works for all distributions

Accuracy (MSE) independent of population size

Parameter Estimation

(+ a Mean Example)



Variance Estimation

(It's Almost Elementary)

Natural estimator

Two calculations

Biased?

Mystery



Estimating the Variance

Unknown distribution or population p

mean μ

variance σ^2

Estimate σ^2

Sample of n observations

$$X_1, \dots, X_n \sim p \perp$$

No distribution

Expectation \rightarrow average

Mean

$$\mu = E(X_i)$$



$$\bar{X} \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n X_i$$

Variance

$$\sigma^2 = E(X_i - \mu)^2$$



$$\text{"S"}^2 \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

"Raw"

Soon

"Raw" sample variance

Random variable

ExSample



Sample

n=5 observations

2, 1, 4, 2, 6

$$\bar{X} = \frac{1}{5} \sum_{i=1}^5 x_i = \frac{2+1+4+2+6}{5} = \frac{15}{5} = 3 \quad \text{Estimated mean}$$

$$\text{"S}^2\text{"} = \frac{1}{5} \sum_{i=1}^5 (x_i - \bar{x})^2 = \frac{1+4+1+1+9}{5} = \frac{16}{5} = 3.2 \quad \text{Estimated variance}$$

$$V(X) = E(X-\mu)^2 = E(X^2)-\mu^2$$

Similar expression for "S²"

One-Pass Calculation

$$\sum_{i=1}^n \rightarrow \sum$$

$$\begin{aligned}\sum(x_i - \bar{x})^2 &= \sum(x_i^2 - 2x_i\bar{x} + \bar{x}^2) \\&= \sum x_i^2 - 2\bar{x} \sum x_i + \sum \bar{x}^2 \\&= \sum x_i^2 - 2n\bar{x}^2 + n\bar{x}^2 & \sum x_i = n\bar{x} \\&= \sum x_i^2 - n\bar{x}^2\end{aligned}$$

$$“S^2”, \stackrel{\text{def}}{=} \frac{1}{n} \sum(x_i - \bar{x})^2$$

Intuitive
Arguments

$$= \frac{1}{n} \sum x_i^2 - \bar{x}^2$$

Fewer subtractions
One pass

ExSample



n=5

2, 1, 4, 2, 6

Saw

$$\bar{X} = 3$$

$$\text{"S}^2\text{"} = 3.2$$

One pass

$$\begin{aligned}\text{"S}^2\text{"} &= \frac{1}{5} \sum_{i=1}^5 x_i^2 - \bar{x}^2 \\ &= \frac{4+1+16+4+36}{5} - 3^2 \\ &= 12.2 - 9 \\ &= 3.2 \quad \checkmark\end{aligned}$$

Even more interesting...

“S²” Biased?

Mean

$$E(X_i) = \mu$$

$E \rightarrow$ average of samples

$$\bar{X} \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n X_i$$

Unbiased

$$E(\bar{X}) = \mu$$

Last lecture
WLLN

Variance

$E \rightarrow$ average of samples

$$E(X_i - \mu)^2 = \sigma^2$$

$$\text{“S}^2\text{”, } \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

Unbiased

?

But is it?

Sherlock Holmes, A study in Scarlet

“No data yet,” he answered.

“It is a capital mistake to theorize before you have all the evidence.

It **biases** the judgment.”

Is “S²” Biased?



Magnifying glass



Simulation

Simulation Plan

Pick a distribution

σ^2 known

Generate n observations

X_1, \dots, X_n

Calculate

$$\text{“}S^2\text{”} = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

Check

$$E(\text{“}S^2\text{”}) \stackrel{?}{=} \sigma^2$$

Find Expectation?

Last lecture, WLLN

$E \approx$ Average of many

r experiments

“ S^2 ” for each

Average $\overline{\text{“}S^2\text{”}}$

$\rightarrow E(\text{“}S^2\text{”})$

Compare

Calculated $\overline{\text{“}S^2\text{”}}$ estimating $E(\text{“}S^2\text{”})$


Known σ^2

Similar



Unbiased

Different



Biased

Normal $N_{0,16}$


$$\text{“}S^2\text{”} = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

n=2

100K Experiments

$$\overline{\text{“}S^2\text{”}} \approx E(\text{“}S^2\text{”})$$

$$E(\text{“}S^2\text{”}) \stackrel{?}{=} \sigma^2$$




Notebook

Exponential E₂

n=2

100K Experiments



Is it bad

Program?

Numbers?

Luck?

Determine exact difference?

Bernoulli



B_p

$P(1) = p$

$P(0) = 1-p = q$

$\sigma^2 = pq$



$n=2$

x_1, x_2

$$\bar{x} = \frac{x_1+x_2}{2}$$

$$\text{“S}^2\text{”}(x_1, x_2) = \frac{1}{2}((x_1 - \bar{x})^2 + (x_2 - \bar{x})^2)$$

x_1, x_2	$P(x_1, x_2)$	\bar{x}	$\text{“S}^2\text{”}$
0,0	q^2	0	$\frac{1}{2} ((0 - 0)^2 + (0 - 0)^2) = 0$
0,1	qp	$\frac{1}{2}$	$\frac{1}{2} \left((0 - \frac{1}{2})^2 + (1 - \frac{1}{2})^2 \right) = \frac{1}{2} \cdot (\frac{1}{4} + \frac{1}{4}) = \frac{1}{4}$
1,0	pq	$\frac{1}{2}$	$\frac{1}{4}$
1,1	p^2	1	0

$$\begin{aligned}
 E(\text{“S}^2\text{”}) &= \sum_{x_1, x_2} p(x_1, x_2) \cdot \text{“S}^2\text{”}(x_1, x_2) \\
 &= q^2 \cdot 0 + qp \cdot \frac{1}{4} + pq \cdot \frac{1}{4} + p^2 \cdot 0 = \frac{pq}{2} = \frac{\sigma^2}{2} !
 \end{aligned}$$

(P)review

$n = 2$

Simulation

$N_{0,16}$ and E_2

$E("S^2")$

$\neq \sigma^2$

$\approx \frac{1}{2} \sigma^2$

Exact calculation

B_p



Indeed $E("S^2") \neq \sigma^2$



$E("S^2") = \frac{1}{2} \sigma^2$ Exactly

Hope!

Other distributions?

Other n ?

Why?

General n

Simulations

$$E("S^2") \approx \frac{n-1}{n} \cdot \sigma^2$$

Variance Estimation

(It's Almost Elementary)

Natural estimator

$$\frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

Two calculations

Two-pass

Bias? Simulation

Analytically

$$E = \frac{n-1}{n} \sigma^2$$



Unbiased
Variance Estimation